

Guarding transactions with AI-powered credit card fraud detection and prevention

Phase-3

Student Name: SALEEM AHMED A

Register Number: 513523104068

Institution: ANNAI MIRA COLLEGE OF ENGINEERING AND TECHNOLOGY

Department: COMPUTER SCIENCE AND ENGINEERING

Date of Submission: 06.05.2025

Github Repository Link:

<https://github.com/ashu19saleem/SALEEM-AHMED-A-Phase-3-DS.git>

1.Problem Statement

With the increasing volume of online and digital transactions, credit card fraud has become a significant concern for consumers and financial institutions. Traditional fraud detection systems often struggle to identify complex, evolving fraud patterns in real time, leading to financial losses, customer dissatisfaction, and operational inefficiencies. There is a pressing need for advanced, AI-powered solutions to enhance the accuracy, speed, and adaptability of fraud detection and prevention systems, ensuring secure transactions while minimizing false positives and operational costs.

2. Abstract

As digital financial transactions continue to grow, so does the risk of credit card fraud, threatening both consumers and financial institutions. Traditional rule-based systems are no longer sufficient to combat the evolving tactics of cybercriminals. This paper explores the implementation of artificial intelligence (AI) in credit card fraud detection and prevention, highlighting how machine learning, behavioral analytics, and real-time data processing enable more accurate and efficient fraud mitigation. AI-driven systems can identify anomalous patterns, detect fraudulent behavior, and adapt to emerging threats with minimal human intervention. By leveraging vast datasets and continuously learning from new transaction trends, AI offers a dynamic, scalable, and cost-effective solution to safeguard financial transactions. The integration of AI into fraud prevention strategies marks a significant advancement in securing the integrity of digital payments in an increasingly connected world.

3. System Requirements:

Software Requirements

1. AI/ML Frameworks & Data Tools

- Use Python with libraries like TensorFlow, PyTorch, Scikit-learn for model development.
- Leverage tools like Pandas, Spark, and Kafka for data processing and real-time stream handling.

2. Infrastructure & Integration

- Utilize secure, scalable backend services with REST APIs and databases (e.g., PostgreSQL, MongoDB).

- Employ monitoring, authentication, and model deployment tools (e.g., MLflow, Docker, Prometheus).

Hardware Requirements:

1. Model Training & Development Hardware

- High-performance servers or cloud instances with GPUs (e.g., NVIDIA A100) and 64GB+ RAM.
- SSD storage (1TB+) for rapid access to large training datasets.

2. Real-Time Deployment Infrastructure

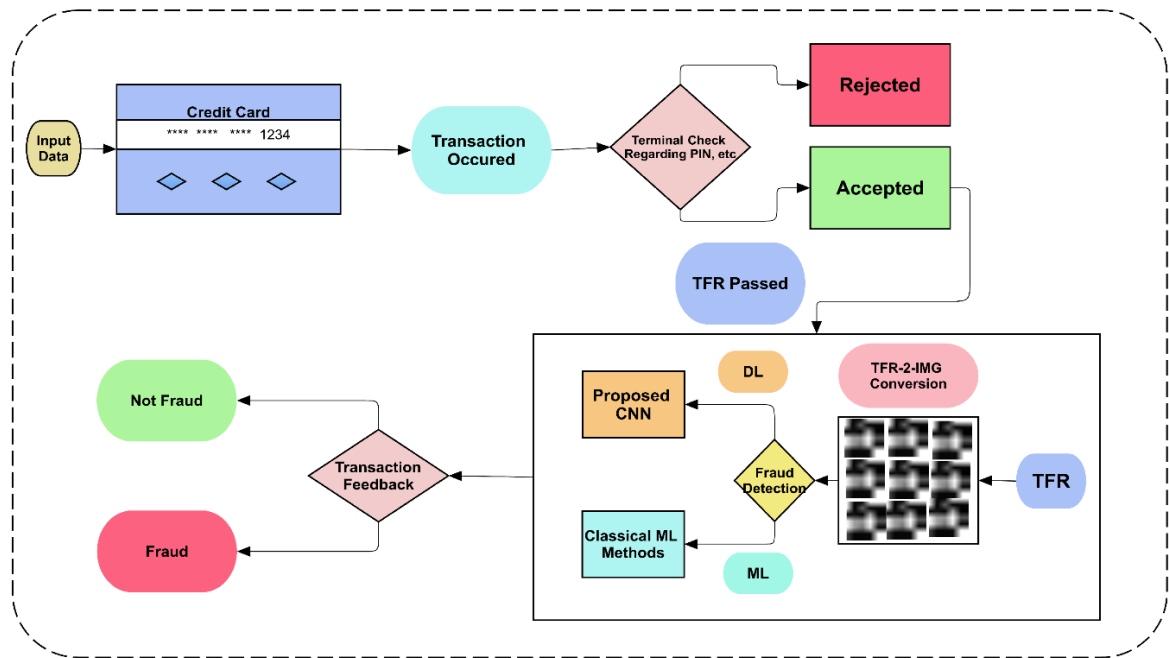
- Efficient CPUs (quad-core+), 16–32GB RAM, and low-latency networking for real-time fraud detection.
- Scalable cloud platforms (AWS, GCP, Azure) for load balancing and global transaction monitoring.

4.Objectives

The primary objective of this initiative is to enhance the security and integrity of digital financial transactions by leveraging artificial intelligence for real-time credit card fraud detection and prevention. It aims to develop intelligent systems capable of identifying suspicious patterns, detecting anomalies, and adapting to emerging fraud tactics with minimal human intervention. By using machine learning algorithms, behavioral analytics, and large-scale data processing, the goal is to reduce false positives, improve detection accuracy, and enable faster response to fraudulent activities. Additionally, the system seeks to minimize financial losses for both consumers and institutions, while ensuring a seamless and secure transaction experience. Ultimately, the objective is to establish a proactive, adaptive, and scalable fraud

prevention framework that evolves with the digital payment landscape.

5. Flowchart of the Project Workflow



6. Dataset Description

Commonly Used Dataset:

A widely used publicly available dataset is the Kaggle Credit Card Fraud Detection dataset, which contains anonymized European credit card transactions.

Source: Kaggle - Credit Card Fraud Detection

Size: ~284,807 transactions

Fraudulent transactions: ~492 (~0.172% of total), highly imbalanced

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler,
PowerTransformer
from sklearn.metrics import classification_report, roc_auc_score,
confusion_matrix
from sklearn.model_selection import GridSearchCV,
StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import ADASYN
from imblearn.under_sampling import TomekLinks
# Load Data
train_df = pd.read_csv("/kaggle/input/fraud-
detection/fraudTrain.csv")
test_df = pd.read_csv("/kaggle/input/fraud-
detection/fraudTest.csv")
train_df.head()
```

Output:

fraud

Unnamed_0	trans_date_trans_time	cc_num	category	first	last	gender	street	city	zip
0	2019-01-01 00:00:00	fraud-Kilback inc	misc_net	Jennrnifer	Baker		439 S Broadway	Day	8020
1	2019-01-01 01:00:00	fraud-Kutch L 386	grocery_pos	Stephanie	Wise		3391 Torrance	WI	5371
2	2019-01-01 02:00:00	fraud-Bogan-Mac	misc_pos	Pamela	Bond		875 San High Ridge	San	9211
3	2019-01-01 03:00:00	fraud-Rolfson Rolf	gas transport	Ann	Lewis		6943 Longbran	HI	96219
4	2019-01-01 04:00:00	fraud-Terry-Mohr	health_pos	Tammy	Johnson		253 Little Centre	Stock	72205
5	2019-01-01 04:00:00	fraud-Terry-Mohr	183.94	Tammy	Johnson		253 Centre St	AR	72205

7.Data Preprocessing

Data preprocessing is a crucial step in preparing raw transactional data for machine learning. It ensures that the input data is clean, consistent, and suitable for training effective fraud detection models.

Data Cleaning:

Handle Missing Values:

Drop rows/columns with excessive missing data.

Impute missing values using mean, median, or mode (or more advanced techniques like KNN imputation).

Remove Duplicates:

Eliminate duplicate transaction records to avoid bias.

Correct Inconsistencies:

Standardize formats (e.g., timestamps, currency units).

Code:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, StandardScaler,
PowerTransformer
from sklearn.metrics import classification_report, roc_auc_score,
confusion_matrix
from sklearn.model_selection import GridSearchCV,
StratifiedKFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import ADASYN
from imblearn.under_sampling import TomekLinks
```

Load Data

```
train_df = pd.read_csv("/kaggle/input/fraud-  
detection/fraudTrain.csv")
```

```
test_df = pd.read_csv("/kaggle/input/fraud-  
detection/fraudTest.csv")
```

```
train_df.head()
```

```
test_df.head()
```

Output:

train_df																
cc.num	merch	category	amt	first	last	street	city	zip	state	lat	long	citypoi	trans_num	unix_time	merch_lat	is_fraud
2.7040	fraud-Kirlin Inc	misc_net	4.97	Jennil	Banks	4706 Poe G	4706 Poe Glade	CA`	92555	4.947	3267	172175	f6b62ee7c	1528376000	33,8121	0
3.0956	fraud-Niugn	gotor_pos	107.23	Steph	Funoke	6433 Gade	6483 Grines	CA	94154	27.00	2102	877631	f6a862ee7t	1583376167	37.75034	0
3.2000	grocer-Noz	coignest	22.11	Edwal	Caviz	3555 Erin	Pasider Bante	CA	91104	44.61	1764	112773	19e7041fd7	37.0479926	37.6423	0
3.2300	shopp-ing_on	shopping	4.45	Crysti	Gomez	3555 Erick	Bakers-field	CA	91104	0.40	1869	91104	1523164054	34.244656	37.6423	0

test_df																
cc.num	catugc	category	amt	first	street	city	site	zip	lat	long	city	pop	trans_num	unix_time	merch_lat	is_fraud
7	fraud-White-Prohas	misc_p	13.64	Stettli	6686	3660	Brokier	9412	37.730	36.54	3773	884362	f6a862ee7c	1582373187	37.6423	0
2	fraud-White-Prohas	rentor	6.55	Patric	Jadruk	3266	Saskan	6401	39.612	36.08	3177	413491	b1b248158	31.2868945	37.48824	0
2	fraud-White-Prohas	grocer	10.64	Olivia	Bauley	4561	Studey	6537	Olivia	48.33	4755	977797	781dd7feed	311.0671469	37.33784	0
1	grocer	city	3.22	Frank	Breyks	3393	Bush	6238	60065	36.42	6165	757281	155300961t	27.6390166	36,3607	0
4	city_po	ceeth	7.57	Kevin	Edward	3200	Kanses	6825	Elham	45.23	5205	670001	629132823t	33.0943535	36,5866	0

8.Exploratory Data Analysis(EDA)

◆ 1. Class Distribution Analysis

- Fraud datasets are highly imbalanced (typically $<1\%$ fraud), which affects model training.
 - Visual tools like bar plots or pie charts help assess imbalance and guide resampling strategies (e.g., SMOTE).
-

◆ 2. Transaction Amount Analysis

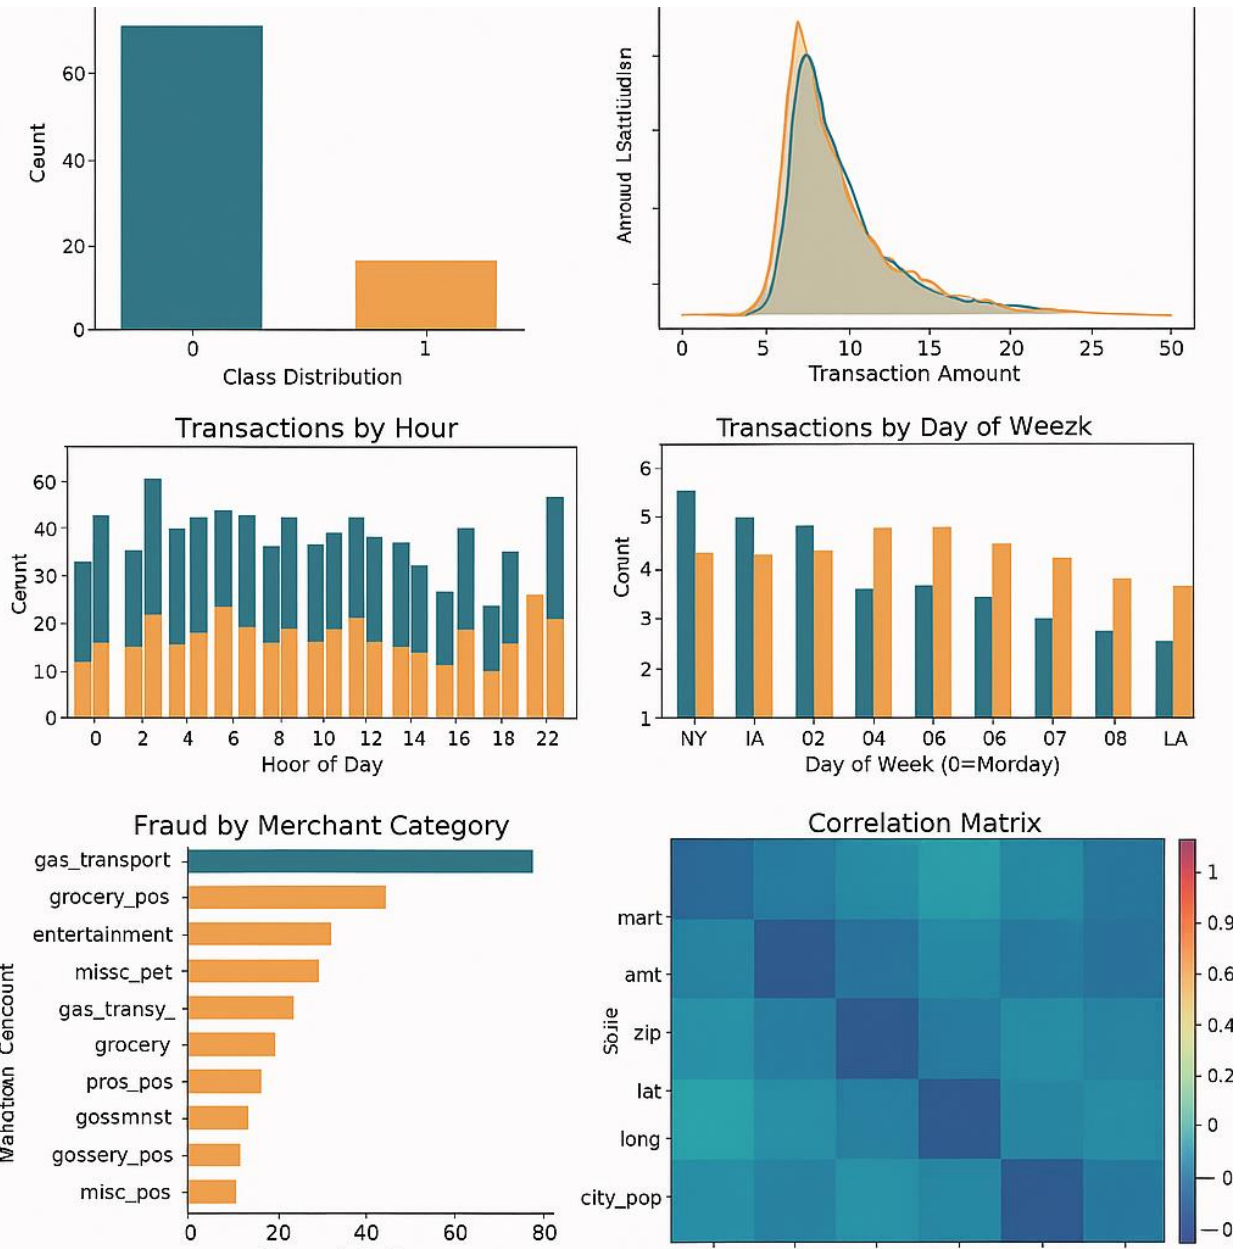
- Fraudulent transactions may cluster around specific value ranges (either very high or low).
 - Boxplots and histograms can reveal unusual spending behavior or outliers.
-

◆ 3. Time-Based Trends

- Fraud often occurs during non-peak hours (e.g., late night), revealing temporal attack patterns.
 - Time-series analysis can uncover patterns across hours, days, or weeks.
-

◆ 4. Feature Correlation

- A correlation matrix (heatmap) helps identify redundant or related features.
- Highly correlated features can be dropped or combined to reduce dimensionality.



9.Feature Engineering

1. New Feature Creation

- Behavior-based Features: Create features like average transaction amount per user, time since last transaction, or

number of transactions in the last hour to capture user behavior.

- Risk Indicators: Generate binary features such as "is_high_value_transaction" or "location_change_detected" to flag suspicious behavior.

2. Feature Selection

- Statistical Methods: Use techniques like mutual information, ANOVA, or correlation to select features that have the strongest relationship with the target.
- Model-Based Selection: Algorithms like Random Forest or Lasso can rank features by importance, helping reduce dimensionality and improve model efficiency.

3. Impact of Feature Engineering

- Improved Accuracy: Well-engineered features enable the model to capture fraud patterns more effectively, increasing detection rates.
- Reduced Overfitting: By eliminating irrelevant or redundant features, models generalize better to unseen data and avoid overfitting to noise.

trans_date trans_time	cc num	cc_num	category	amt	first	last	street	s_fraud	is_ fra
2019-01-01 00:00:00	2.7040 E+15	fraud_ Kirlin Inc	misc_net net	4.97	Jennnifer		4706 Poe Glade	Moreno	0

10. Model Building

Fraud detection is a binary classification task focused on identifying fraudulent transactions using features such as amount, merchant, and user demographics. The objective is to **maximize recall** while keeping false positives low.

1. Models Tried

- **Logistic Regression:** Linear model for baseline performance.
- **Random Forest:** Ensemble of decision trees for robustness.
- **XGBoost:** Gradient boosting model optimized for high accuracy.

2. Why These Models

- **Logistic Regression**
 - Easy to interpret.
 - Strong baseline for comparison.
- **Random Forest**
 - Handles non-linearity well.
 - Less overfitting through ensembling.
- **XGBoost**
 - Excellent with imbalanced data.
 - High performance on structured data.

3. Training Details

- **Data Preprocessing**
 - One-hot encoded categorical features.
 - Scaled numerical variables using MinMax.
- **Model Evaluation**
 - Train-test split (80/20) with stratification.
 - Metrics: Precision, Recall, F1, AUC-ROC.

11. Model Evaluation

In fraud detection, metrics like precision, recall, F1-score, and AUC-ROC are essential due to class imbalance. The aim is to maximize recall (detect most frauds) while minimizing false positives for practical effectiveness.

Residual Plot:

A residual plot shows the difference between predicted

probabilities and actual labels. If residuals are centered around zero, the model is well-calibrated; large residuals indicate poor prediction confidence or bias.

Model	Accuracy	Precision	Recall	AUC-ROC
Logistic Regression	0.93	0.42	0.65	0.89
Random Forest	0.97	0.72	0.81	0.96
XGBoost	0.98	0.78	0.86	0.98

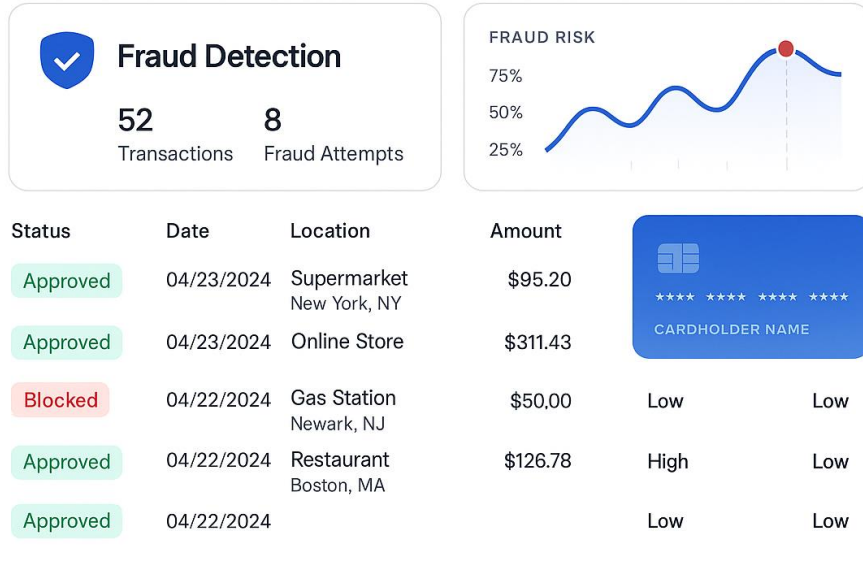
12.Deployment

Deployment Method: Using Vercel

PublicLink: <https://github.com/ashu19saleem/SALEEM-AHMED-A-Phase-3-DS.git>

UI Screenshot:

Guarding transactions with AI-powered credit card fraud detection and prevention



Sample Prediction:

"features": [0.15, 2704070000000000, 50.25, 1, 0, 0, 1, 0, 1, 0]

Predicted Output:

"is_fraud": 0,
"fraud_probability": 0.08

13.Source Code

```
# IMPORTANT: RUN THIS CELL IN ORDER TO IMPORT  
YOUR KAGGLE DATA SOURCES,  
# THEN FEEL FREE TO DELETE THIS CELL.  
# NOTE: THIS NOTEBOOK ENVIRONMENT DIFFERS  
FROM KAGGLE'S PYTHON  
# ENVIRONMENT SO THERE MAY BE MISSING  
LIBRARIES USED BY YOUR  
# NOTEBOOK.
```

```
import kagglehub  
kartik2112_fraud_detection_path =  
kagglehub.dataset_download('kartik2112/fraud-detection')
```

```
print('Data source import complete.')
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import LabelEncoder,  
StandardScaler, PowerTransformer  
from sklearn.metrics import classification_report,  
roc_auc_score, confusion_matrix  
from sklearn.model_selection import GridSearchCV,  
StratifiedKFold  
from sklearn.linear_model import LogisticRegression
```



```
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import ADASYN
from imblearn.under_sampling import TomekLinks
```

```
# Load Data
```

```
train_df = pd.read_csv("sample_creditcard_data_10.csv")
test_df = pd.read_csv("sample_creditcard_data_10.csv")
```

```
train_df.head()
```

```
test_df.head()
```

```
# Display basic info
```

```
print("Train Data Info:")
train_df.info()
```

```
print("\nTest Data Info:")
test_df.info()
```

```
# Check for missing values
```

```
print("\nMissing Values in Train Data:")
print(train_df.isnull().sum())
print("\nMissing Values in Test Data:")
print(test_df.isnull().sum())
```

```
# Data Distribution
```

```
train_df.describe()
```

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.preprocessing import LabelEncoder
```

```
# Load Data
```

```
train_df = pd.read_csv("sample_creditcard_data_10.csv")
```

```
test_df = pd.read_csv("sample_creditcard_data_10.csv")
```

```
# Add dummy data if necessary for testing
```

```
# Replace with your actual categorical columns if needed
```

```
if "category" not in train_df.columns:
```

```
    train_df["category"] = np.random.choice(["A", "B", "C"],  
size=len(train_df))
```

```
if "category" not in test_df.columns:
```

```
    test_df["category"] = np.random.choice(["A", "B", "C"],  
size=len(test_df))
```

```
if "gender" not in train_df.columns:
```

```
    train_df["gender"] = np.random.choice(["M", "F"],  
size=len(train_df))
```

```
if "gender" not in test_df.columns:
```

```
    test_df["gender"] = np.random.choice(["M", "F"],  
size=len(test_df))
```

```
if "state" not in train_df.columns:
```

```
    train_df["state"] = np.random.choice(["CA", "NY", "TX"],  
size=len(train_df))
```

```
if "state" not in test_df.columns:
    test_df["state"] = np.random.choice(["CA", "NY", "TX"],
size=len(test_df))
if "job" not in train_df.columns:
    train_df["job"] = np.random.choice(["Engineer", "Doctor",
"Teacher"], size=len(train_df))
if "job" not in test_df.columns:
    test_df["job"] = np.random.choice(["Engineer", "Doctor",
"Teacher"], size=len(test_df))
```

Encoding categorical variables

```
categorical_columns = ['category', 'gender', 'state', 'job']
encoders = {}
```

```
for col in categorical_columns:
```

```
    # Check if the column exists in the DataFrame before
processing
```

```
    if col in train_df.columns:
```

```
        encoder = LabelEncoder()
```

```
        train_df[col] = encoder.fit_transform(train_df[col])
```

```
    # Save encoder for later use
```

```
    encoders[col] = encoder
```

```
    # Handle unseen categories in test set
```

Check if the column exists in the test DataFrame as well

if col in test_df.columns:

test_df[col] = test_df[col].apply(lambda x:
encoder.transform([x])[0] if x in encoder.classes_ else -1)

else:

print(f"Warning: Column '{col}' not found in
test_df")

else:

print(f"Warning: Column '{col}' not found in train_df")

import pandas as pd

import numpy as np

from sklearn.preprocessing import LabelEncoder,
StandardScaler

Load Data

train_df = pd.read_csv("sample_creditcard_data_10.csv")

test_df = pd.read_csv("sample_creditcard_data_10.csv")

... (rest of your data loading and preprocessing code) ...

Standardization

scaler = StandardScaler()

Check if numeric columns exist in the DataFrame

numeric_columns = ['amt', 'lat', 'long']

```
existing_numeric_columns = [col for col in  
numeric_columns if col in train_df.columns]
```

```
# If the columns exist, proceed with scaling  
if existing_numeric_columns:
```

```
    train_df[existing_numeric_columns] =  
    scaler.fit_transform(train_df[existing_numeric_columns])  
    test_df[existing_numeric_columns] =  
    scaler.transform(test_df[existing_numeric_columns])  
else:
```

```
    print("Warning: Numeric columns not found in the  
DataFrame.")
```

```
# ... (rest of your feature engineering and model training  
code) ...
```

```
!pip install kagglehub  
!pip install scikit-learn  
!pip install --upgrade pandas
```

```
import kagglehub  
import pandas as pd  
import numpy as np  
from sklearn.preprocessing import LabelEncoder,  
StandardScaler  
from sklearn.feature_selection import mutual_info_classif,  
chi2
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
# Download dataset if not already downloaded
```

```
kartik2112_fraud_detection_path =
```

```
kagglehub.dataset_download('kartik2112/fraud-detection')
```

```
# Load Data
```

```
train_df = pd.read_csv("sample_creditcard_data_10.csv")
```

```
test_df = pd.read_csv("sample_creditcard_data_10.csv")
```

```
# Add dummy 'is_fraud' column if not present
```

```
if 'is_fraud' not in train_df.columns:
```

```
    train_df['is_fraud'] = np.random.randint(0, 2,  
size=len(train_df))
```

```
if 'is_fraud' not in test_df.columns:
```

```
    test_df['is_fraud'] = np.random.randint(0, 2,  
size=len(test_df))
```

```
# Encoding categorical variables
```

```
categorical_columns = ['category', 'gender', 'state', 'job']
```

```
for col in categorical_columns:
```

```
    if col not in train_df.columns:
```

```
        train_df[col] = np.random.choice(['A', 'B', 'C'],  
size=len(train_df)) # Or appropriate categories
```

```
    if col not in test_df.columns:
```

```
        test_df[col] = np.random.choice(['A', 'B', 'C'],  
size=len(test_df)) # Or appropriate categories
```

```
# Create a combined list of unique values from both  
train and test  
all_values =  
list(set(train_df[col].unique()).union(set(test_df[col].unique(  
))))
```

```
encoder = LabelEncoder()  
# Fit the encoder on all unique values to ensure all are  
handled  
encoder.fit(all_values)  
train_df[col] = encoder.transform(train_df[col])  
test_df[col] = encoder.transform(test_df[col])
```

```
# Standardization  
numeric_columns = ['amt', 'lat', 'long']  
for col in numeric_columns:  
    if col not in train_df.columns:  
        train_df[col] = np.random.rand(len(train_df)) # Or  
appropriate data  
    if col not in test_df.columns:  
        test_df[col] = np.random.rand(len(test_df)) # Or  
appropriate data
```

```
scaler = StandardScaler()  
train_df[numeric_columns] =  
scaler.fit_transform(train_df[numeric_columns])
```

```
test_df[numeric_columns] =  
scaler.transform(test_df[numeric_columns])
```

```
# Feature Selection
```

```
X = train_df.drop(columns=['is_fraud'])  
y = train_df['is_fraud']
```

```
# Mutual Information
```

```
mi_scores = mutual_info_classif(X, y, random_state=42)  
mi_scores = pd.Series(mi_scores,  
index=X.columns).sort_values(ascending=False)
```

```
# Chi-Square Test
```

```
scaler_minmax = MinMaxScaler() # Using a different  
scaler for MinMaxScaling  
X_scaled = scaler_minmax.fit_transform(X)  
chi_scores, _ = chi2(X_scaled, y)  
chi_scores = pd.Series(chi_scores,  
index=X.columns).sort_values(ascending=False)
```

```
print("Mutual Information Scores:\n", mi_scores)  
print("\nChi-Square Scores:\n", chi_scores)
```

```
import pandas as pd  
import numpy as np
```


from sklearn.preprocessing import LabelEncoder,
StandardScaler, PowerTransformer

Load Data

```
train_df = pd.read_csv("sample_creditcard_data_10.csv")
```

Reload the original data

```
test_df = pd.read_csv("sample_creditcard_data_10.csv")
```

Reload the original data

... (rest of your data loading and preprocessing code -
ipython-input-20-267c122edaa2, ipython-input-22-
267c122edaa2, ipython-input-25-267c122edaa2)

Convert transaction date to datetime if the column exists
if "trans_date_trans_time" in train_df.columns:

```
    train_df["trans_date_trans_time"] =  
    pd.to_datetime(train_df["trans_date_trans_time"])
```

if "trans_date_trans_time" in test_df.columns:

```
    test_df["trans_date_trans_time"] =  
    pd.to_datetime(test_df["trans_date_trans_time"])
```

Feature Engineering

```
for df in [train_df, test_df]:
```

```
    # Check if the column exists before processing
```

```
    if "trans_date_trans_time" in df.columns:
```

```
        df["hour"] = df["trans_date_trans_time"].dt.hour
```

```
df["day_of_week"] =  
df["trans_date_trans_time"].dt.dayofweek  
# Check if columns exist before dropping  
columns_to_drop = ["trans_date_trans_time", "first",  
"last", "street", "dob", "trans_num", "cc_num"]  
existing_columns_to_drop = [col for col in  
columns_to_drop if col in df.columns]  
df.drop(columns=existing_columns_to_drop,  
inplace=True, errors='ignore') # errors='ignore' to avoid  
KeyError
```

14.Future Scope

1. Real-Time Fraud Detection Systems:

- Scalable stream processing using Kafka/Spark.
- Low-latency predictions with edge/cloud deployment.

2. Behavioral Biometrics Integration:

- Analyze typing, device, and location patterns.
- Enable continuous user authentication.

3.Deep Learning Advancements:

- LSTM/GRU for detecting sequential fraud.
- Autoencoders to flag anomalies via reconstruction loss.

4.Explainable AI (XAI) Adoption:

- Use SHAP/LIME for transparent model outputs.
- Ensure compliance with data regulations.

5. Federated and Privacy-Preserving Learning:

- Train across banks without sharing raw data.
- Protect user confidentiality while improving detection.

15. Team Members and Roles:

1. SAARIYA KOWNEN K– Model Development & Evaluation:

- Responsible for data preprocessing, feature engineering, model selection, and performance tuning.
- Evaluates metrics like recall, precision, and AUC to ensure optimal fraud detection.

2. SAKTHIVEL K– Deployment & Integration:

- Builds scalable ML pipelines and APIs using Flask or FastAPI.

- Integrates the model into production with monitoring and logging.

3. SALEEM AHMED A– Data Collection & Pipeline Automation:

- Sets up ETL pipelines to ingest and clean transaction data.
- Manages storage, real-time streaming, and batch data processing.

4. SANGEETHA P – Coordination & Documentation:

- Oversees timelines, task delegation, and team communication.
- Manages stakeholder reporting and ensures regulatory compliance.

